

AquaSecure

MSP · Operator Console

Functional Documentation

A screen-by-screen walkthrough of the provider-side console — what each menu does and what the operator can do.

This reference covers the AquaSecure provider console below, then continues into the customer-facing **Tenant Console** in Part 2. Most provider screens are multi-tenant by default — they either aggregate the whole fleet or break the same view down per tenant.


1. Analysis

The day-to-day operational surface — real-time situational awareness across the fleet, the consolidated alert queue, ingestion health and platform observability.

1.1 Command Center

The default landing screen after login. It aggregates the entire fleet into one pane, organised into tabs.


Overview



Command Center — Overview: fleet-wide KPI strip, alert and event analytics, and per-tenant storage tables.

A live, fleet-wide snapshot. The KPI strip tracks headline numbers — active tenants, new events, alerts and detections. Analytics cover API ingestion over time and alerts by tenant and by severity (critical / high / medium / low). Bottom tables show **Federated TIP** (per-provider health, quota and IOC cache) and per-tenant **Customer storage**.


Threat Detection



Command Center — Threat Detection: live global attack map with verdict breakdown and a streaming detections feed.

A live global threat map. The header summarises the current window — attack rate, total detections and verdict mix. Detections are plotted worldwide and colour-coded Malicious / Suspicious / Benign / Unknown, with a streaming Recent Detections feed (time, source country, IP, tenant, verdict).


AI Detection



Command Center — AI Detection: MSP-only cross-tenant IOC observations from Aqua-TIP / Aqua-Hunter, with an IP detail panel.

An MSP-only, cross-tenant view of IOC observations produced by **Aqua-TIP** and **Aqua-Hunter**. A detections table (time, agent, tenant, IOC, verdict with confidence) pairs with an IP-detail panel — cache verdict, first / last seen, evidence and GEO.


Threat Intel



Command Center — Threat Intel: the consolidated AquaCTI knowledge base — reports, indicators, actors, malware and campaigns.

The consolidated **AquaCTI** knowledge base behind the fleet. KPI tiles count Reports, Indicators, Observables, Threat Actors, Malware, Vulnerabilities, Campaigns and Intrusion Sets. Report ingestion is trended over 30 days and split by type (Threat-Report, News-Report, Ransomware-Report, MISP-Event), alongside a knowledge-base-by-type breakdown, recently added threat actors, and “most active” vulnerability / threat / malware lists.


Analyst Activity



Command Center — Analyst Activity: real-time in-app usage for MSP analysts and tenant users, with a per-user session table.

A real-time view of who is actually using the console — in-app activity for both MSP analysts and tenant users in one pane. KPI cards show Active now (last 2 min), Active last 24h, total time-in-app, and analyst / tenant-user counts. A stacked time-in-app chart (5-minute buckets) sits above a per-user table: kind and role, tenant, active time, sessions, tenants touched, switches, actions and last seen.


1.2 Alerts Central



Alerts Central: the consolidated cross-tenant alert queue with summary donuts, a timeline, filters and the alert table.

The consolidated, cross-tenant alert queue. A time-range selector and search sit above six summary donuts and a timeline. A filter rail (tenant, severity, status ...) narrows the alert table — event time, tenant, title, category, nature, severity, status, assignee, matches, and an Open action.


1.3 API Monitoring



API Monitoring: ingestion and storage pipeline health, with per-tenant ingest gauges and source health.

The health of the data pipeline. KPI cards track API ingestion (events/sec), ETL and DB ingestion. Per-tenant ingest gauges pair a trailing-window ring with a rate ring; a health row covers hourly volume (ingested / rejected / errors), storage history by tenant, and per-source status.

1.4 Observability



Observability — Overview: platform / infra health, the OpenTelemetry service map, and per-subsystem metrics.

Platform observability for AquaSecure’s own infrastructure, built on **OpenTelemetry**. The Overview maps services and per-subsystem metrics; further tabs add a live Logs explorer, distributed Traces with a span tree, a Metrics catalogue, and Storage & Retention for the observability store.

2. Access

Governs threat-intel sourcing, the operator audit record, and storage across tenants.

2.1 Audit Trail

MSP Audit Trail
4779 events — append-only log

Keyword search... All events All time Filter by operator email...

Satisfies ISO 27001:2022 A.5.18 (Access rights) and A.5.15 (Logging).

Timestamp	OPERATOR	EVENT	DETAIL	IP
6/22/2026, 8:28:58 PM		Login		
6/22/2026, 8:28:49 PM		Login Password Ok		
6/22/2026, 8:28:43 PM		Impersonation Action		
6/22/2026, 8:27:36 PM		Impersonation Action		
6/22/2026, 8:26:12 PM		Impersonation Action		
6/22/2026, 8:24:43 PM		Impersonation Action		
6/22/2026, 8:23:36 PM		Token refresh		
6/22/2026, 8:23:36 PM		Impersonation Action		
6/22/2026, 8:23:36 PM		Login		
6/22/2026, 8:22:12 PM		Token refresh		
6/22/2026, 8:22:12 PM		Impersonation Action		
6/22/2026, 8:20:43 PM		Impersonation Action		
6/22/2026, 8:19:36 PM		Impersonation Action		
6/22/2026, 8:18:12 PM		Impersonation Action		
6/22/2026, 8:16:43 PM		Impersonation Action		
6/22/2026, 8:15:36 PM		Impersonation Action		
6/22/2026, 8:14:12 PM		Impersonation Action		
6/22/2026, 8:12:43 PM		Impersonation Action		
6/22/2026, 8:11:36 PM		Impersonation Action		
6/22/2026, 8:11:02 PM		Token refresh		
6/22/2026, 8:10:13 PM		Impersonation Action		
6/22/2026, 8:08:43 PM		Impersonation Action		
6/22/2026, 8:07:36 PM		Impersonation Action		
6/22/2026, 8:06:12 PM		Impersonation Action		
6/22/2026, 8:04:43 PM		Impersonation Action		


Audit Trail: the append-only MSP operator activity log, mapped to ISO 27001 controls.

An append-only record of operator activity, explicitly mapped to compliance (ISO 27001:2022 A.5.18 and related controls). Columns cover timestamp, operator, event, detail and source IP; the event vocabulary includes login (TOTP / email), password OK, MFA email sent, token refresh, and the tenant-impersonation events.

3. Integration

How external data gets into AquaSecure and how it is normalised into the platform's canonical event shape.

3.1 SIEM Sources



The screenshot shows the AquaSecure Operator Console interface. On the left is a dark sidebar with navigation menus for MSP, Analysis, Manage, Access, API, Content, AI Agents, and Tools. The main panel is titled 'SIEM Sources' and contains a table of configured ingest endpoints per tenant. The table has columns: TENANT, VENDOR, MODE, LABEL, LAST ACTIVITY, RATE LIMIT, STATUS, and ACTIONS. Two sources are listed:

TENANT	VENDOR	MODE	LABEL	LAST ACTIVITY	RATE LIMIT	STATUS	ACTIONS
Wazuh (mppl edge)	Dmarc	Pull	shared	polled 4m ago	600/min	active	Pull now, Edit, Setup, Revoke, Delete
Wazuh (mppl edge)	Wazuh	Push	Wazuh (mppl edge)	—	5000/min	active	Test, Edit, Setup, Revoke, Delete

The 'Push' mode source is highlighted with a blue box. Below the table, there is a note: 'Configured ingest endpoints per tenant. Push-mode receives events from the vendor; pull-mode polls the vendor's REST API on a schedule.'

SIEM Sources: per-tenant ingest endpoints in push or pull mode, with rate limits and status.

Per-tenant ingest endpoints. The defining choice is the ingest mode — **Push** (the vendor sends events to AquaSecure) or **Pull**. Adding a source selects the tenant and vendor (DMARC Aggregate, Phishbot, IBM QRadar, Splunk, Wazuh ...), with rate limits and status per source.


4. Content

The GRC authoring studio — branded **AquaPilot**. Its guiding model is author once, publish to many tenants.

4.1 GRC Design

A multi-tab studio for authoring and curating compliance frameworks, controls, mappings, questionnaires and monitors.

Frameworks



GRC Design — Frameworks: the framework library the MSP authors and publishes.

The framework library the MSP authors and publishes. KPI cards track Active (tenant-visible), Draft (MSP-only) and total adoptions; the catalogue spans regulators and standards.

Framework detail


The screenshot displays the AquaSecure Operator Console interface. On the left is a dark sidebar with navigation menus for Access, Integration, Content, Agents, and Tools. The main panel shows the 'GRC Design' section with a breadcrumb trail: Frameworks > Controls > Sub-elements > Cross Mapping > Maturity Mapping > Maturity Assessment > Tree View > Gap Assessments > Monitor Catalog > Tenant Access > Version Control. The selected framework is 'Critical Telecom Data and Infrastructure Security Regulations' (v2020 - PK - PTA - parent: 15627861), which is active. It has 23 criteria, 68 mappings, 2 templates, and 1 monitors. Below this, a table lists 'Controls & sub-elements' with columns for ID, name, risk score, and actions. The table includes entries for 'SE-GR-037' (Information security events involving suppliers are managed and tracked through to resolution, risk 80) and 'SE-GR-032' (Information security requirements are addressed in agreements with suppliers and reviewed before signature, risk 70). A detailed view of 'SE-GR-037' shows cross-mappings to various frameworks (ISO 27001, NIST CSF, CIS, SOC, GAO, etc.) and a list of document templates (SOP-TPR-001, AGR-SLA-001). The bottom of the page shows a list of other frameworks like CTDSR-AWARE, CTDSR-BCM, CTDSR-CHANGE, etc.

GRC Design — Framework detail: the control hierarchy with risk, mappings, monitors and document templates.

Opening a framework reveals its full control hierarchy. Each control → sub-element carries a risk score, automated monitors, cross-mappings and document templates — the canonical model that powers compliance-satisfy-many.

5. AI Agents


5.1 SOC Agent Studio



SOC Agent Studio: the autonomous SOC roster — L2 specialists, the L3 hunter, and the L1 router.

The low-code builder for the autonomous SOC — build, configure and deploy agents by dragging blocks onto a canvas. The roster runs **L0–L3**: L0 (Aqua_L0, pre-hook) dedupes every incoming alert; L1 (aqua_l1) is a deterministic keyword / IOC router with no LLM; L2 specialists (network, identity, endpoint, webapp, email, cloud, malware) investigate; and L3 (aqua_l3) is a scheduled threat hunt sweeping raw events. Supporting agents include Aqua-TIP (IOC validation) and Aqua-Hunter (scheduled SIEM sweep).

5.2 AI Grid



AI Grid — Alert Triage Pipeline: the live left-to-right view of every alert flowing from ingest through detection and routing.

A single-screen visualisation of the autonomous SOC — the same agents shown as a live, left-to-right pipeline. Stages run Ingest (the aqua-proxy connector and per-tenant SIEM / DMARC sources feed the msp-api gateway) through detection, routing and L2 / L3 handling.

6. Account

6.1 My Security

The screenshot displays the 'My Security' page in the AquaSecure Operator Console. The left sidebar contains a navigation menu with categories: MSP (Analysis, Manage, Access, API, Content, AI Agents, Tools), and My security. The main content area is titled 'My security' and includes a 'Security settings' section with three main components:

- Two-factor authentication:** A section for managing 2FA. It shows 'Authenticator app' is configured. Below, it states 'You have an authenticator app configured. When you sign in, you'll be asked for the 6-digit code from your app.' There are buttons for 'Regenerate backup codes' and 'Disable authenticator'.
- Password:** A section for changing the account password. It includes fields for 'CURRENT PASSWORD', 'NEW PASSWORD', and 'CONFIRM NEW PASSWORD'. A note specifies 'Min 12 characters: 3 of 4: uppercase, lowercase, digit, symbol'. An 'Update password' button is at the bottom.
- Active sessions:** A section showing devices currently signed in. It includes a 'Sign out all' button and a table of active sessions.

DEVICE	IP	STARTED	
Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7) A...	181.53.248.19	6/9/2026, 8:52:53 PM	Revoke
Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7) A...	181.53.248.19	6/9/2026, 6:59:44 PM	Revoke

My Security: per-operator account security — 2FA, password policy, and active sessions.

The operator's personal account-security panel — distinct from org-level Access. It covers two-factor authentication (authenticator-app TOTP, with backup-code regeneration), password (change under a min-12-character, 3-of-4-class policy), and active sessions (every signed-in device, with per-session revoke and sign-out-everywhere).

AquaSecure

Tenant Console

Functional Documentation

A screen-by-screen walkthrough of the customer-facing console — the security workspace each client organisation logs into.


Active tenant context: multinet.mppl.aquasecure.ai · Part 2 of 2 · Revision July 2026

This is the tenant end of the product. The provider end — the **MSP · Operator Console** the security team runs — is Part 1 above. Modules are grouped into Core (Governance & Compliance, Security Ops, Digital Exposure, Threat Intelligence) and Essentials (Operations, Assets & Identity, Security Tools, and more).

1. vCISO Command Center

The tenant's single-pane situational view — live posture, alert analytics and current SOC activity.


1.1 Overview



vCISO Command Center — Overview: live posture KPIs, alert and event analytics, noise reduction, and current SOC activity.

The default landing screen, rolling the whole organisation's posture into one time-scoped view. A KPI strip tracks GRC readiness, cyber maturity, open cases, critical / high alerts, wallet balance and storage. Analytics cover alert-severity trends, incoming data and EPS, and breakdowns by source, user, category and status. A **Noise Reduction** panel shows how much raw volume the L0 layer absorbed; **Who's Currently Working** tracks live AI-agent, automation and human activity; and the bottom row surfaces DMARC, the Phishbot funnel, asset health and a live threat feed.

1.2 Alerts




vCISO Command Center — Alerts: breakdown donuts, an events-in-window timeline, and the filterable alert table.

A focused view of the tenant’s own alert queue. Summary donuts break the window down by source, severity, status, category and assignee; an events timeline is colour-coded by severity; and a filter rail + table list event time, title, category, nature, severity, status and assignee.

2. Security Ops

The tenant's hands-on detection surface — searching raw and alert data, writing correlation rules, and querying an external SIEM.


Dashboard



Security Ops — Dashboards: the SOC Overview board — events, severity, top hosts, MITRE techniques and open incidents at a glance.

A pre-built SOC analytics workspace. The **SOC Overview** tab summarises events (raw), high + critical volume, active sources and distinct hosts, with an event-volume-by-severity chart, severity distribution, top hosts, top MITRE techniques and a recent high-severity events table. Further tabs cover Authentication Activity, MITRE ATT&CK Coverage, Vulnerabilities, Network Flow, CIS Benchmark, Alerts & Cases Ops and Asset & Vuln Posture — with a Builder for custom boards.

2.1 Search




Security Ops — Search: the SIEM search surface with pinned facets, an alert-rows / raw-events histogram, and the results table.

Free-form search across the tenant's security data. A query bar (field:value syntax) with time-range chips sits above an alert-rows / raw-events histogram and pinned facets (severity, source, category, device class, host, actor / target user). A results table lists time, severity, source, host, category, device class / vendor, title and count.

3. Digital Exposure

The tenant's outward-facing view — what an attacker can see from the internet. It contains External Attack Surface discovery, DMARC monitoring and Phishbot.


3.1 Overview



Digital Exposure — Overview: exposure score, findings-by-severity, top exposed assets, and live scan activity.

The exposure posture at a glance, scoped to the external footprint over a 30-day window. A KPI strip shows Exposure Score (lower is better), active findings (with a critical count), seeds monitored and scans. An exposure trend plots score against finding count; findings break down by severity; top exposed assets (including look-alike phishing domains) are ranked by finding count; and scan activity shows discovery seeds by status.


3.2 Attack Surface



Digital Exposure — Attack Surface: discovery KPIs, status and scan-type donuts, and the discovery-seed register.

External attack-surface discovery. The tenant registers discovery seeds — domains, IPs or CIDR ranges — and each is scanned. KPI and donut panels track total seeds, scan status (completed / running / pending / failed) and the scan-type mix (domain, IP-range, SSL/TLS, email security, dark-web). A seed register lists every seed with type, status, findings and re-scan controls.

3.3 DMARC



Digital Exposure — DMARC: alignment KPIs, per-domain status, and the live GRC posture tie-in.


Email-authentication monitoring across the tenant's domains over a 7 / 30 / 90-day window. A KPI strip covers domains monitored, aligned %, failing % and reporting organisations. Per-domain status shows

records seen, alignment %, SPF / DKIM pass %, and reject / quarantine counts. A live **GRC posture** panel evaluates SPF / DKIM / DMARC and refreshes evidence against the Email Security control.

3.4 Phishbot

The tenant's phishing-simulation and security-awareness module — design lure emails and fake landing pages, target groups, deliver campaigns and measure results.

Dashboard




Phishbot — Dashboard: the aggregated campaign funnel and per-recipient results.

Aggregated results across all campaigns — email sent, opened, clicked and submitted as a funnel — plus a per-recipient breakdown. Build blocks (groups, email templates, landing pages), delivery (sending profiles, campaigns) and results all hang off this view.

Threat Intelligence

The tenant's window onto **AquaCTI** — the platform's global threat-intelligence graph. It spans a live threat map, a searchable knowledge base, dark-web monitoring, supply-chain vendor risk, a CVE tracker, phishing intelligence and a curated news feed.

Threat Map



Threat Intelligence — Threat Map: a live world map of current attacks, with recent indicators, latest reports and active actors.

A live world map of current attacks. Panels summarise global threats (latest attacks, countries, attack types), top source countries and a live feed, with recent indicators (IPv4, URL, X.509), latest threat reports and active threat actors down the side.

Explore

The screenshot shows the AquaCTI Explore interface. On the left is a sidebar with navigation options: Marketplace, Connectors, All Modules, vCISO Command Center, CORE (Governance & Compliance, Security Ops, Digital Exposure, Threat Intelligence, Vendor Risk, CVE Tracker, PhishWatch, Threat News), ESSENTIALS (Operations, Assets & Identity, Knowledge, Security Tools), and Employee portal. The main panel is titled 'Explore' and contains a search bar, a filter bar with categories like Reports, Indicators, Observables, Malware, Threat Actors, Intrusion Sets, Campaigns, Vulnerabilities, Tools, and Attack Patterns, and a search bar for entities. Below the search bar, there are 9,603,018 results. The results are displayed as a list of entities, each with a TLP marking and a timestamp. The entities include indicators, hostnames, domain names, and IP addresses.

Threat Intelligence — Explore: search the AquaCTI knowledge base — indicators, malware, actors, reports and more.

A direct search into the AquaCTI knowledge base — indicators, malware, actors, reports and more. Entity-type tiles (Reports, Indicators, Observables, Malware, Threat Actors, Intrusion Sets, Campaigns, Vulnerabilities, Tools, Attack Patterns) filter a unified results list, each row carrying a TLP marking and timestamp.


Dark Web

The screenshot shows the AquaCTI Dark Web interface. On the left is a sidebar with navigation options: Marketplace, Connectors, All Modules, vCISO Command Center, CORE (Governance & Compliance, Security Ops, Digital Exposure, Threat Intelligence, Vendor Risk, CVE Tracker, PhishWatch, Threat News), ESSENTIALS (Operations, Assets & Identity, Knowledge, Security Tools), and Employee portal. The main panel is titled 'Dark Web' and contains a search bar, a filter bar with categories like Social Vectors, Live Feed, Top Darknets, Target Vectors, and Top Vectors, and a search bar for entities. Below the search bar, there are 688,728 results. The results are displayed as a list of entities, each with a TLP marking and a timestamp. The entities include indicators, hostnames, domain names, and IP addresses.

Threat Intelligence — Dark Web: live dark-web listings, leaks and breach chatter across a large index.

Live dark-web listings, leaks and breach chatter across a large index. Social-vector and top-darknet panels sit beside a live feed (source, date, content, author, class), with target vectors by country and a ranked breakdown of leak categories — combo lists, stealer logs, ransomware, database dumps, credit cards and more.


Vendor Risk



Threat Intelligence — Vendor Risk: supply-chain intelligence, ranking affected vendors by recent CVE volume.

Supply-chain intelligence from AquaCTI. The **Vendors at Risk** tab ranks affected vendors by recent CVE volume over a selectable window, each with a per-vendor CVE count and last-seen time.


Intel Browser



Vendor Risk — Intel Browser: a searchable CVE browser with aggregated risk and severity breakdown.

The second Vendor Risk tab — a searchable CVE browser. Rows list severity, CVE, source (NIST NVD, Shodan ...), description and detection date, alongside an aggregated-risk gauge and a severity breakdown (critical / high / medium / low, with a known-exploited count).


CVE Tracker



Threat Intelligence — CVE Tracker: the live CVE landscape from AquaCTI (CISA KEV · NVD · OSV).

The live CVE landscape from AquaCTI (CISA KEV · NVD · OSV). KPI cards total CVEs, known-exploited (KEV) and each severity band. Severity and CVSS-score distributions, a by-source split, a CVE trend line, top affected vendors (by severity) and a top-critical-CVEs table complete the picture.

PhishWatch



Threat Intelligence — PhishWatch: live phishing intelligence — URLs, domains, IPs and files with public references.

Live phishing intelligence from AquaCTI — URLs, domains, IPs and files with public references. KPI cards count phishing IOCs, URLs, domains, files, IPs and attack techniques. A phishing-IOC volume trend, IOC-type mix, top phishing types and top feeds / sources sit above a live feed (IOC, type, score, source, first-seen, and public references such as urlscan and VirusTotal).

Threat News

The screenshot shows the AquaSecure Tenant Console interface. On the left is a sidebar with navigation options: Marketplace, Connectors, and Threat Intelligence. The main content area is titled "Threat News" and displays "Latest threat intelligence reports - newest first". At the top right of the main area is a tag treemap showing various categories like security, hacking, information, security n..., artificial, privacy, vulner..., and malware. Below the treemap is a search bar and filters for categories and sources. The main content area lists several threat intelligence reports, each with a title, source, timestamp, and a summary. The reports include:

- From Phishing to Persistence: A CrySome RAT Infection Chain Analysis** by AlienVault, Jul 7, 2026 7:14 PM, 366 linked. Summary: A sophisticated multi-stage infection chain was analyzed following successful containment by MDR SOC operations. Initial access occurred through spear-phishing using a logistics rate confirmation lure, delivering CrySome remote access trojan via multiple stages. The attack chain leveraged living-off-the-land techniques, ICMUaUtil COM interface for UAC bypass, and in-memory AMSI patching. WinDefChl, an open-source Defender disruption tool... Tags: smi bypass, browser hijacking, credential theft, +5 more. TLP: CLEAR. Read source.
- qilin has published a new victim: Accelerate** by Ransomware.Live, Jul 7, 2026 6:59 PM, 8 linked. Summary: N/A. TLP: CLEAR. Read source.
- Fake Netflix, Coca-Cola, and FIFA job scams target marketers** by AquaSecure, Jul 7, 2026 6:43 PM. Summary: A fake recruiter phishing campaign uses trusted brands, nested redirects, and fake Google prompts to steal accounts. Tags: google credentials, job offers, news, +2 more. Read source.
- Court Filing Reveals Windows Device ID Helped FBI Trace Alleged Scattered Spider Hacker** by AquaSecure, Jul 7, 2026 6:27 PM. Summary: U.S. prosecutors linked an alleged Scattered Spider hacker to a break-in at a luxury jewelry retailer using a persistent Windows device ID, according to a newly unsealed federal complaint. Microsoft records tied that ID first to the account the attackers used to keep access during the May 2025 intrusion, then to online accounts prosecutors say belong to 19-year-old Peter Stokes. Stokes is...


Threat Intelligence — Threat News: the latest threat-intelligence reports, newest first.

The latest threat-intelligence reports, newest first. A tag treemap surfaces trending topics; a searchable, filterable feed lists each report with source, timestamp, linked-entity count, a summary, tags and a TLP marking.

4. Operations

The day-to-day SOC workspace — the alert queue and its full investigation workflow, cases and agent automation.


4.1 Alerts



Operations — Alerts: the tenant's SOC alert queue with filters, assignees and per-row actions.

The tenant's full alert queue. Filters cover search, category, nature, status, severity, priority, source, coverage, assignee and a date range; the table lists event time, title (with a match-count badge), category, source, nature, severity, status and assignee (AI or human).


Alert details



Operations — Alerts → Alert details: the full record, AI-analysis actions, workflow controls, notes and linked events.

Opening an alert reveals everything about it — severity, priority, source and event / first-seen / last-seen times — plus AI-analysis actions and workflow controls (create case, create rule, notify), notes, and linked activity, contributing events and aggregated child alerts.


View AI analysis



Alert details — View AI analysis: the structured aiAnalysis record already produced by the agent pipeline.

Shows the structured analysis the AI agent already produced for the alert — routing (which L2 agent and model), verdict, severity recommendation and the reasoning behind it — the aiAnalysis record generated by the agent pipeline.

4.2 AgentFlow



Operations — AgentFlow: the tenant's AI-agent and scheduled-flow automation, wired to its own alert pipeline.

Where the tenant turns on AI agents and scheduled flows against its own alert stream. Slot-usage cards track capacity; new flows are schedule-driven rules with the same filter shape as a workflow rule; and the agent catalog exposes the same **L0-L3** roster the provider manages centrally, here as enable-and-bill options priced per month in credits.

5. Assets & Identity

The tenant's inventory of what it has to protect and who is responsible for it.

5.1 Assets Dashboard


The screenshot shows the AquaSecure Tenant Console interface. The left sidebar contains navigation options: Security Ops, Dashboards, Search, BYOD, Security Workflow, Digital Exposure, Threat Intelligence, ESSENTIALS, Operations, Statistics, Alerts, Cases, AgentFlow, Notifications, Assets & Identity (selected), Knowledge, Security Tools, and Settings. The 'Assets & Identity' section is expanded, showing 'Assets Dashboard' as the active view. The main content area displays the 'Asset list (49 of 49)' with a search bar and filters for All, Wazuh agents, Syslog devices, and Unknown. The table lists assets with columns for Hostname, Type, Status, OS, Version, IP, Vendor, CVE-C, CVE-H, Sub-tenant, and Last seen. The table shows a mix of 'Wazuh agent' and 'Syslog device' types, with statuses ranging from 'Online' to 'Offline'. The 'Last seen' column indicates when each asset was last detected, ranging from 'just now' to 'Jun 19'.

Hostname	Type	Status	OS	Version	IP	Vendor	CVE-C	CVE-H	Sub-tenant	Last seen
192.168.1.10	Wazuh agent	Online	—	—	192.168.1.10	Wazuh	—	—	—	just now
192.168.1.11	Syslog device	Online	—	—	192.168.1.11	Wazuh	—	—	—	just now
192.168.1.12	Wazuh agent	Online	—	—	192.168.1.12	Wazuh	—	—	—	2m ago
192.168.1.13	Wazuh agent	Online	—	—	192.168.1.13	Wazuh	—	—	—	2m ago
192.168.1.14	Wazuh agent	Online	—	—	192.168.1.14	Wazuh	—	—	—	3m ago
192.168.1.15	Wazuh agent	Offline	—	—	192.168.1.15	Wazuh	—	—	—	4d ago
192.168.1.16	Wazuh agent	Offline	—	—	192.168.1.16	Wazuh	—	—	—	4d ago
192.168.1.17	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 20
192.168.1.18	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.19	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.20	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.21	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.22	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.23	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.24	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.25	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.26	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.27	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.28	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.29	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19
192.168.1.30	Unknown	Offline	—	—	—	Wazuh	—	—	—	Jun 19

Assets & Identity — Assets Dashboard: the discovered-asset inventory across agents, syslog devices and unknowns.

The live inventory of everything seen on the network. Filter by hostname, IP, OS, vendor or sub-tenant; each row shows hostname, type, online / offline status, OS and version, IP, vendor, critical / high CVE counts and the owning sub-tenant.


6. Security Tools



Security Tools: the toolkit landing grid — client-side utilities and DNS-record lookups.

A built-in toolkit of client-side utilities and DNS-record lookups — quick checks that never leave the browser. Utilities include a password analyzer (strength, entropy, crack-time), hash generator (SHA-256 / 512 / 1) and CIDR calculator; email / DNS lookups resolve and validate DMARC, SPF and DKIM records, plus a DMARC generator.

7. AI Assistance



AI Assistance: the console's conversational security copilot, reachable from the top bar on any screen.


The console's conversational security copilot, reachable from the top bar on any screen. A chat list holds past conversations — including auto-created ALERT threads (the per-alert write-ups) — while the conversation pane renders full markdown analyses (severity, headline, actors, IOC references, scope). Ask anything about security, or about a specific alert.

Sovereign AI. AI Assistance runs on AquaSecure's sovereign-AI stack — the models are hosted inside the platform's own infrastructure, so tenant data and security context stay within the deployment rather than being sent to a third-party AI provider.

Employee Portal

A self-service space each staff member logs into — separate from the SOC console — reachable from the tenant sidebar. It surfaces the security work assigned to that person (policies, training, campaigns and tasks) and lets them install and check the endpoint agent on their own device. A **Viewing as** switcher at the top previews the portal as any mock employee identity.

Overview



Employee Portal — Overview: a personal greeting, Open / Overdue / Completed counters, and outstanding policies and training.

The employee's landing screen. A personal greeting is followed by three headline counters — Open, Overdue and Completed — then the outstanding items grouped into **Policies to acknowledge** and **Training to complete**, each card showing type, status, due date and a completion badge.

My Assignments

The screenshot displays the 'My Assignments' page in the AquaSecure Employee Portal. The interface includes a sidebar with navigation options: Overview, My Assignments (selected), History, and Endpoint. The main content area is titled 'VIEWING AS Senior Software Engineer' and contains several sections:

- Policies to acknowledge:** A single item, 'Annual Acceptable Use Policy attestation 2026', is marked as 'ACKNOWLEDGED'.
- Training to complete:** This section lists enrolled courses with progress bars and completion status:
 - Cybersecurity Awareness Fundamentals:** 100% complete, Score: 92%, Completed 2026-03-16.
 - Phishing Attack Recognition:** 100% complete, Score: 88%, Completed 2026-03-22.
 - GDPR & Data Privacy Essentials:** 65% complete, 'Continue learning' button.
 - Cloud Security Best Practices (AWS):** 30% complete, 'Continue learning' button.
- ASSIGNED CAMPAIGNS:** Two items are listed as 'COMPLETED':
 - Phishing 101 refresh — Q2 2026 (2024-07-07).
 - Secure Coding — Engineering onboarding (2024-07-06).
- Assigned tasks & incidents:** A section for tasks and incidents, currently showing a task titled 'Upload evidence for CE-1700 — Quarterly Access Review'.

Employee Portal — My Assignments: policies, enrolled courses with progress, assigned campaigns, and tasks & incidents.

The full to-do list for that employee. **Policies to acknowledge** sits above **Training to complete** — enrolled courses with a percent-complete bar, score and a Continue-learning action — then **Assigned campaigns** (phishing-awareness and onboarding tracks) and **Assigned tasks & incidents** (GRC tasks such as evidence uploads, each with a type and priority).

History


The screenshot displays the 'History' section of the AquaSecure Employee Portal. The interface includes a top navigation bar with a 'Menu' icon, a 'STANDARD' plan indicator, '5,000' users, and a 'VIEWING AS Senior Software Engineer' dropdown. The left sidebar shows navigation options: 'Overview', 'My Assignments', 'History' (selected), and 'Endpoint'. The main content area displays a table of completed training, acknowledged policies, and finished campaigns.

Item	Date	Status
✓ Phishing 101 refresh - Q2 2026	2026-07-07	ACKNOWLEDGED
✓ Phishing 101 refresh - Q2 2026	2026-07-07	COMPLETED
✓ Secure Coding - Engineering onboarding	2026-07-05	COMPLETED

Employee Portal — History: a time-stamped record of completed training, acknowledged policies and finished campaigns.

A record of everything the employee has already cleared — completed training, acknowledged policies and finished campaigns — each row time-stamped with its final status (Completed / Acknowledged).

Endpoint



Employee Portal — Endpoint: self-service install of the AquaSecure Endpoint Agent, with a full monitoring scope and a privacy note.

The self-service endpoint-agent screen. It shows install status for the **AquaSecure Endpoint Agent** and a one-click install, above a “What the agent monitors” grid — Access Control & Authentication, Configuration & Hardening, Vulnerability Management, Data & File Integrity, Asset & Software Inventory, Threat Detection & Malware, Audit Log Management, and Incident & Response Activity.

Privacy. The screen is explicit that the agent collects only security-relevant system events — not personal files, browser history, private messages, emails, or camera / microphone data. Only those events are reported to the organisation’s security team.